
Leant

a Djex-based synthesis REPL for Lean 4

an overview and tutorial, with a detailed tour of `:synth`

An **experimental** tool in active development. Interfaces, output, and capabilities change frequently.
github.com/VladimirReshetnikov/Leant August 2026.

Contents

1	Introduction	2
2	Getting started	2
3	The commands	3
3.1	Session and evaluation	3
3.2	Discovery	4
3.3	Proving and synthesis	4
3.4	Persistence and utility	4
4	Interactive proving	5
5	Term synthesis: <code>:synth</code>	8
5.1	Higher-order plumbing	9
5.2	Programs you already know	10
5.3	Rank-N and impredicative goals	11
5.4	Impossibility, proved	17
5.5	Inductive types	17
5.6	Recursion from the library	19
5.7	Classical candidates	20
5.8	Dependent formulas as cargo	21
5.9	Synthesis inside a proof	21
5.10	Engines and budgets	23
5.11	What <code>:synth</code> will not do	29
6	How it works, briefly	30
7	Status	31

1 Introduction

Leant is an interactive read-eval-print loop for Lean 4, in the spirit of GHCi: you type expressions and they are evaluated, you type declarations and they enter the session, and a family of colon-commands gives you type queries, documentation, search, interactive proving, and automatic *term synthesis* with machine-checked answers. The synthesis command takes up half of this document.

Lean 4 itself is normally driven from an editor: the language server shows goals and diagnostics as you edit a file. Leant complements that workflow with a conversational one. It is aimed at exploration (trying a lemma, poking at a definition, asking *is this type even inhabited?*), where the unit of work is a line, not a file.

experimental

Leant is an **experimental project in active development**. It is a research vehicle, not a supported product: commands appear, change shape, and disappear between commits; output formats are not stable; and some features are bounded approximations by design. Every transcript in this document was produced by the tool at the date on the title page and may not match the tool you build tomorrow. When this document and `:help` disagree, believe `:help`.

Leant is implemented in Haskell and speaks the backend protocol directly; term synthesis (§5) embeds the Djex synthesis library in-process.

Under the hood, Leant drives the community `leanprover-community/repl` binary over a JSON protocol: every input becomes a command against a persistent Lean environment, and every claim Leant makes, including every synthesized term, is checked by that backend before it is shown. Sessions survive backend crashes (the session history replays automatically), can be saved and restored (`:pickle/:unpickle`, including synthesis-ready companions for Leant-created snapshots), and can be recorded as transcripts.

2 Getting started

Build the binary once with `cabal build exe:leant` (GHC 9.12.4; the vendored `lib/Djex` submodule must be initialized), then run `leant.cmd`, or the built executable directly. Leant finds the Lean backend the same way `LeanInteract`'s cache laid it out, or accepts an explicit `--repl-exe/LEANT_BACKEND`. Without a Lake project it runs against the plain Lean 4 standard library, which is also the least memorable thing about it: no imports, subsecond startup.

The command line

session

```
leant [FILE] [--project DIR] [--plain] [-i MOD] [--timeout N] [--time]
      [--transcript [FILE]] [--timestamps] [--repl-exe PATH] [--lake PATH]
      [--length-ranking-config ABSOLUTE-PATH]
      [--length-ranking-config-timeout MS] [--length-ranking-allow-unpinned]
```

A positional FILE is loaded as by `:load`. `--project DIR` runs inside a Lake project (auto-detected from the working directory without the flag), and `--plain` forces the bare-stdlib session shown above. `-i MOD` (also `--import`, repeatable, comma-splitting) adds startup imports; `--timeout N` sets the per-command budget in seconds (0 disables it, default 300); `--time` shows per-command timing; `--transcript [FILE]` and `--timestamps` start transcript recording; `--lake PATH` selects the lake executable. The three `--length-ranking` options activate finite-list-spine Length behavioral assessment (§5.10): the configuration file is admitted once at startup from an explicitly chosen absolute path, the timeout bounds only the file-load interruption budget (default 5,000 ms, at most 60,000 ms), and both auxiliary options are rejected without the configuration option itself. The current startup file is versionless and requires `rankingDomain` to be exactly `scalar` or `binary-product`. The separate one-command contract file is also one current versionless shape: exactly `format`, `rankingDomain`, and `contract`. The selected nested contract requires explicit target roles and case policy, admits the current modulo and positive-literal quotient grammar, and has no `resultShape` field.

A session begins as a calculator and escalates quietly:

```

session
λ> 2 + 2
4
λ> it * 10
40
λ> def double (n : Nat) : Nat := n + n
λ> double 21
42
λ> :type double
double : Nat → Nat

```

Expressions are tried with `#eval` and fall back to `#check`; the last evaluated value is bound as `it`, GHCi-style. Declarations (`def`, `theorem`, `inductive`, ...) extend the session environment, and any `#-`command (`#eval`, `#check`, `#print axioms`) passes straight through. Multi-line input opens automatically when a line is syntactically incomplete (type a `structure` header and just keep going), and an empty line submits it; an explicit `{ ... }` block delimits one by hand, as in GHCi. TAB completes `:` commands and dotted identifiers against the live environment.

3 The commands

3.1 Session and evaluation

<code>:help, :h, :?</code>	show the command summary
<code>:quit, :q</code>	exit
<code>:type EXPR, :t</code>	show the type of EXPR (<code>#check</code>)
<code>:info NAME, :i</code>	show the definition of NAME (<code>#print</code> , rendered as a valid declaration)
<code>:load FILE, :l</code>	reset the session and load a <code>.lean</code> file
<code>:reload, :r</code>	reload the last loaded file
<code>:import MOD</code>	add an import (rebuilds the session; unavailable over an opaque snapshot)

<code>:imports</code>	list active imports, or imports restored by <code>:reset</code>
<code>:undo</code>	revert the last state-changing command
<code>:reset</code>	clear definitions or a snapshot base (keeps configured imports)
<code>:history</code>	show commands after the current import/snapshot base
<code>:set OPT VAL</code>	a Lean <code>set_option</code> that persists in the session

3.2 Discovery

<code>:browse [NS]</code>	list declarations in a namespace, or the session's own
<code>:browse! NS</code>	...including compiler-generated auxiliaries
<code>:doc NAME</code>	show a docstring
<code>:search TEXT</code>	search declaration names, case-insensitively
<code>:search? TYPE</code>	proof search: what proves TYPE? (via <code>exact?</code>)

3.3 Proving and synthesis

<code>:prove [PROP]</code>	prove PROP interactively, tactic by tactic (without an argument: resume the last sorry)
<code>:synth TYPE</code>	synthesize verified terms of TYPE; candidates are bound as <code>it1 (= it)</code> , <code>it2</code> , ... (§5)
behavior option	<code>--behavior-mode rank filter</code> explicitly chooses the operation; disabled rank is identity, while filtering uses the activated startup Length contract and requires its policy
contract option	<code>--length-contract PATH</code> supplies one passive Length contract file for this command only; an optional preceding behavior mode selects rank or filter, omitted mode means rank, and every contract form requires the startup policy opt-in (§5.10)
<code>:set synth-engine E</code>	<code>djinn</code> (default) <code>exference</code> <code>both</code>
<code>:set synth-steps N</code>	Exference's step budget (default 4096)
<code>:set synth-classical B</code>	offer classical candidates for constructively refuted goals (<code>on off</code> , default <code>on</code>)
<code>:set synth-library B</code>	library premises for recursive inductives (<code>on off</code> , default <code>on</code> ; §5.6)
<code>:set synth-timeout N</code>	seconds one <code>:synth</code> may search before answering "no answer" (default 20, or <code>LEANT_SYNTH_TIMEOUT</code> at startup; 0 waits indefinitely)

Exact option-bearing grammar

```
:synth [--behavior-mode rank|filter] [--length-contract PATH] -- TYPE
```

3.4 Persistence and utility

<code>:transcript [F on off]</code>	record a transcript of the session
<code>:timestamps [on off]</code>	timestamp transcript entries
<code>:pickle FILE</code>	save the environment and synthesis companion
<code>:unpickle FILE</code>	restore an environment as a new undo/history base
<code>:env</code>	show the backend environment id

```
:time          toggle per-command timing
:! CMD         run a shell command
```

4 Interactive proving

`:prove` turns the prompt into a tactic loop against the backend's proof-state protocol. Every line is a tactic; `:undo [N]` takes back the last N tactics (default one, without limit); `:goals` reprints the goals; `:script` shows the accumulated proof; `:auto` tries a battery of common finishers; `:abort` leaves the mode without saving; and `:qed NAME` assembles the accumulated script into a theorem and replays it into the session, so the result is usable afterwards. A `?-tactic` such as `exact?` records the tactic it *found* rather than the question-mark form. Each displayed proof state is followed by an automatically discovered tactic suggestion that Lean has already checked. The candidate probes are shaped by the goal and its hypotheses: an `intro` suggestion names the binders it would introduce, a disjunction goal is probed with `left/right`, a hypothesis whose type a single step can take apart is probed with `cases` or `obtain`, and a data-typed variable the goal mentions is probed with `induction`. The search prefers a candidate that closes the goal outright and annotates the suggestion accordingly (closes the goal, splits into 2 goals); when no single tactic closes it, a second phase chains quick finishers onto the best progressing candidates (constructor `<=>` `omega`, or an `obtain` step chained with the `exact` it enables), so even the opening suggestion is often a complete checked proof. The chains also try `simp_all`, unfolding the definitions the goal mentions and calling in `omega` on whatever arithmetic remains, which is how an induction suggestion can arrive as a finished proof: `induction l <=> simp_all [myLen]`, or `induction n <=> simp_all [double] <=> omega`. The suggestion is advisory—it does not advance the state or enter the script—and `:suggest` reprints it without rerunning the search. Below, the very first suggestion is already a complete proof; the session ignores it and proceeds step by step instead. Proof-state identifiers are local to one backend process; if it stops, Lean leaves `prove` mode and prints the accumulated script rather than submitting a stale identifier after session reconstruction.

session

```
λ> :prove ∀ p q : Prop, (p → q) → p → q ∧ p
entering prove mode – type tactics; :help for commands
⊢ ∀ (p q : Prop), (p → q) → p → q ∧ p
suggestion: exact fun p q a a_1 => (a a_1, a_1) (closes the goal)
⊢> intro p q h hp
p q : Prop
h : p → q
hp : p
⊢ q ∧ p
suggestion: exact (h hp, hp) (closes the goal)
⊢> exact (h hp, hp)
All goals accomplished
finish with :qed [NAME], inspect with :script
⊢> :qed mp_and
saved: theorem mp_and : ∀ p q : Prop, (p → q) → p → q ∧ p
```

Accepting a suggestion is just typing it. On a conjunction of two arithmetic facts, the opening suggestion is already the whole chained proof, and the session is two lines long:

session

```
λ> :prove (∀ a b : Nat, a + b = b + a) ∧ (∀ a b : Nat, a + 2 * b = b + b + a)
entering prove mode – type tactics; :help for commands
⊢ (∀ (a b : Nat), a + b = b + a) ∧ ∀ (a b : Nat), a + 2 * b = b + b + a
suggestion: constructor <;> omega (closes the goal)
⊢> constructor <;> omega
All goals accomplished
finish with :qed [NAME], inspect with :script
⊢> :qed two_omegas
saved: theorem two_omegas : (∀ a b : Nat, a + b = b + a) ∧ (∀ a b : Nat, a + 2 * b = b
↪ + b + a)
```

That chain was suggested because *both* halves fall to omega. When they need different tools, the suggestion adapts to whichever goal is in focus — arithmetic for the first half, a found library lemma for the second:

session

```
λ> :prove (∀ a b : Nat, a + b = b + a) ∧ (∀ a b : Nat, a * b = b * a)
entering prove mode – type tactics; :help for commands
⊢ (∀ (a b : Nat), a + b = b + a) ∧ ∀ (a b : Nat), a * b = b * a
suggestion: constructor (splits into 2 goals)
⊢> constructor
– goal 1 of 2 –
case left
⊢ ∀ (a b : Nat), a + b = b + a
– goal 2 of 2 –
case right
⊢ ∀ (a b : Nat), a * b = b * a
suggestion: omega (closes the goal)
⊢> omega
case right
⊢ ∀ (a b : Nat), a * b = b * a
suggestion: exact fun a b => Nat.mul_comm a b (closes the goal)
⊢> exact fun a b => Nat.mul_comm a b
All goals accomplished
finish with :qed [NAME], inspect with :script
⊢> :qed add_mul_comm
saved: theorem add_mul_comm : (∀ a b : Nat, a + b = b + a) ∧ (∀ a b : Nat, a * b = b *
↪ a)
```

The probes are shaped by the hypotheses as much as by the goal. An existential hypothesis draws an obtain suggestion that names its witness; once the conjunction inside is the thing in the way, the suggested step is the obtain that splits it, chained with the exact it enables:

session

```
λ> :prove (∃ x : Nat, 2 < x ∧ x < 4) → ∃ y : Nat, 2 < y
entering prove mode – type tactics; :help for commands
⊢ (∃ x, 2 < x ∧ x < 4) → ∃ y, 2 < y
suggestion: intro h
```

```

h> intro h
h :  $\exists x, 2 < x \wedge x < 4$ 
h>  $\exists y, 2 < y$ 
suggestion: obtain (x, h1) := h
h> obtain (x, h1) := h
x : Nat
h1 :  $2 < x \wedge x < 4$ 
h>  $\exists y, 2 < y$ 
suggestion: obtain (h, h2) := h1 <;> exact Exists.intro x h (closes the goal)
h> obtain (h, h2) := h1 <;> exact Exists.intro x h
All goals accomplished
finish with :qed [NAME], inspect with :script
h> :qed ex_two_lt
saved: theorem ex_two_lt :  $(\exists x : \text{Nat}, 2 < x \wedge x < 4) \rightarrow \exists y : \text{Nat}, 2 < y$ 

```

A suggestion is not limited to one step. Here the goal needs structural induction on a user-defined function; once `intro` brings the variable into the context, the suggested chain arrives as a complete verified proof:

session

```

λ> def myLen : List Nat → Nat
...> | [] => 0
...> | _ :: t => myLen t + 1
...>
λ> :prove  $\forall l : \text{List Nat}, \text{myLen } l = l.\text{length}$ 
entering prove mode – type tactics; :help for commands
h>  $\forall (l : \text{List Nat}), \text{myLen } l = l.\text{length}$ 
suggestion: intro l
h> intro l
l : List Nat
h> myLen l = l.length
suggestion: induction l <;> simp_all [myLen] (closes the goal)
h> induction l <;> simp_all [myLen]
All goals accomplished
finish with :qed [NAME], inspect with :script
h> :qed myLen_length
saved: theorem myLen_length :  $\forall l : \text{List Nat}, \text{myLen } l = l.\text{length}$ 

```

The same pattern scales to arithmetic recursion. The chain below was assembled from three of the probes at once: `induction` because the goal mentions a data-typed variable, `simp_all` unfolding the definition the goal is about, and `omega` for the arithmetic the simplifier leaves behind – a finished proof of a small theorem about a function defined moments earlier:

session

```

λ> def double : Nat → Nat
...> | 0 => 0
...> | n + 1 => double n + 2
...>
λ> :prove  $\forall n : \text{Nat}, \text{double } n = 2 * n$ 
entering prove mode – type tactics; :help for commands
h>  $\forall (n : \text{Nat}), \text{double } n = 2 * n$ 
suggestion: intro n

```



```

-> intro n
n : Nat
- double n = 2 * n
suggestion: induction n <=> simp_all [double] <=> omega (closes the goal)
-> induction n <=> simp_all [double] <=> omega
All goals accomplished
finish with :qed [NAME], inspect with :script
-> :qed double_two_mul
saved: theorem double_two_mul : ∀ n : Nat, double n = 2 * n

```

A sorry in an ordinary declaration prints its goal and offers the same loop via bare :prove. Prove mode also cooperates with :synth; §5.9 shows what that buys.

5 Term synthesis: :synth

:synth TYPE answers the question “*write me a term of this type*”; read through propositions-as-types, that is “*prove this*.” Sometimes it answers the stronger question “*show me that no such term exists*.” It is built on the vendored Djex library, linked in-process. Djex began by merging two classic Haskell program synthesizers behind one checked synthesis foundation: **Djinn**, a complete and terminating proof search for intuitionistic propositional logic (Dyckhoff’s LJ_T calculus) that emits programs, and **Exference**, a ranked heuristic search under explicit budgets. It has since moved well beyond both originals. It fixes a long list of their bugs; the worst of them let a search that had merely approximated a type report that type uninhabited, and an exhausted search over an approximated space now counts for nothing. It implements much better type-class support, with a unified, validated constraint contract shared by both engines: Exference resolves givens, superclasses, and instances, while Djinn checks contexts and synthesizes dictionary-independent terms. And it adds a growing level of support for rank-*N* and impredicative types: goal-side quantifiers open through polarity-aware plan families, quantified hypotheses are instantiated at a bounded set of sequent-supplied types, a separate final tail admits closed forall-free monotypes from the checked query, and retained loaded schemes keep their own environment-backed family. Impredicative instantiation is admitted under a guard that never invents a polytype the query did not supply. Djinn’s current context-free hypothesis rule reaches six leading binders without raising its fixed search caps; chains with seven or more stay opaque. Its singleton, pairwise, triple, quadruple, and bounded quintuple open/opaque frontiers cover every choice through eleven independent quantified sites without enumerating a power set. The quintic categories retain at most 512 plans per orientation; a twelve-site six-open/six-opaque choice is the next deliberate occurrence gap. That scope is still being improved, and Leant keeps the boundary to the engines narrow on purpose, so each improvement arrives by version bump. The engines speak Haskell types; Leant translates Lean goals into their fragment and their answers back into Lean.

Three rules run through the design:

- **The engine is never trusted.** Every candidate is re-elaborated by the Lean backend against the exact goal (example : (T) := term) before you see it. A rendering bug or an engine bug costs a dropped candidate, never a wrong answer.

- **Refusals come with reasons.** A goal outside the supported fragment is turned away with a note saying what fell outside, not quietly mangled into something answerable.
- **Negative verdicts are labeled by strength.** “Provably uninhabited” appears only when the translation was complete and lossless; anything weaker is reported as “no term found within bounds,” which claims nothing.

Transcripts in this section are lightly abridged: trailing candidates beyond the ones shown, and trailing truncation notes, are elided, as is the preparing synthesis environment ... line that precedes a session’s first :synth. The celebration emoji that follows All goals accomplished is omitted for the benefit of the document fonts. Every other line that does appear is a verbatim record of a live session.

5.1 Higher-order plumbing

The sweet spot is the structural fragment: arrows, products, sums, $\perp$, $\top$, negation, Iff, and quantifiers over opaque variables. These are the “plumbing” terms one writes constantly. Free capital identifiers are auto-bound, so quick queries stay quick:

```
session
λ> :synth ((A → B → C) → (A → B) → A → C)
  it1 fun f g x => f x (g x)
λ> :synth (A × (B ⊕ C) → (A × B) ⊕ (A × C))
  it1 fun ⟨x, y⟩ => match y with | .inl z => .inl ⟨x, z⟩ | .inr w => .inr ⟨x, w⟩
λ> :synth (∀ p q : Prop, (p ↔ q) → ¬p → ¬q)
  it1 fun _ _ (⟨_, f⟩) k x => k (f x)
```

The first is the S combinator; the second distributes a product over a sum; the third is contraposition across an Iff, with the backward implication projected out by the pattern $\langle _, f \rangle$. Binders are named by role (functions f g h , values x y z , negations and continuations k), a small thing that makes large candidates much easier to read.

The everyday combinators all come out on the first try, and each first candidate is the term you would have written yourself — flip, composition, curry, uncurry, and product associativity:

```
session
λ> :synth ((a → b → c) → b → a → c)
  it1 fun f x y => f y x
λ> :synth ((b → c) → (a → b) → a → c)
  it1 fun f g x => f (g x)
λ> :synth ((A × B → C) → A → B → C)
  it1 fun f x y => f ⟨x, y⟩
λ> :synth ((A → B → C) → A × B → C)
  it1 fun f ⟨x, y⟩ => f x y
λ> :synth (((A × B) × C) → A × (B × C))
  it1 fun ⟨⟨x, y⟩, z⟩ => ⟨x, ⟨y, z⟩⟩
```

Candidates are ranked smallest-first and *bound into the session* as it1, it2, ..., with bare it the best one. They are ordinary definitions; you can evaluate them immediately:

session

```

λ> :synth (a → a → a)
  it1 fun _ x => x
  it2 fun x _ => x
λ> #eval it2 "left" "right"
"left"

```

Both inhabitants of $a \rightarrow a \rightarrow a$ that matter, and picking one is just using its name. For programs, as opposed to proof-irrelevant propositions, the choice between candidates is the whole point.

An Iff goal is treated as a pair of implications, so both directions of a logical identity arrive in one anonymous constructor. De Morgan and the absorption law:

session

```

λ> :synth (∀ p q : Prop, ¬(p ∨ q) ↔ ¬p ∧ ¬q)
  it1 fun _ _ => {fun k => {fun x => k (.inl x), fun y => k (.inr y)}, fun {k1, k2} z
    ↪ => match z with | .inl w => k1 w | .inr x1 => k2 x1}
λ> :synth (∀ p q : Prop, p ∧ (p ∨ q) ↔ p)
  it1 fun _ _ => {fun {x, _} => x, fun y => {y, .inl y}}

```

Under propositions-as-types, the proof terms *are* plumbing, and the synthesized spellings make that visible: Iff symmetry is a swap, and $\neg(p \wedge q) \leftrightarrow (p \rightarrow \neg q)$ is nothing but currying:

session

```

λ> :synth (∀ p q : Prop, (p ↔ q) → (q ↔ p))
  it1 fun _ _ {f, g} => {g, f}
λ> :synth (∀ p q : Prop, ¬(p ∧ q) ↔ (p → ¬q))
  it1 fun _ _ => {fun k x y => k {x, y}, fun k1 {z, w} => k1 z w}

```

The textbook curiosities are in range too. No proposition is equivalent to its own negation, and the synthesized proof is the classic self-application trick — the term below applies k to itself through f twice, the propositional shadow of $\omega \omega$:

session

```

λ> :synth (∀ p : Prop, (p ↔ ¬p) → False)
  it1 fun _ {k, f} => k (f (fun x => k x x)) (f (fun y => k y y))

```

5.2 Programs you already know

Some types have one sensible inhabitant, and asking for it by type is quicker than remembering which library corner it lives in. The bind of the reader monad, then the bind of the state monad:

session

```

λ> :synth ((S → A) → (A → S → B) → S → B)
  it1 fun f g x => g (f x) x

```

```

λ> :synth ((S → A × S) → (A → S → B × S) → S → B × S)
it1 fun f g x => match f x with | ⟨y, z⟩ => g y z
it2 fun f g x => match f x with | ⟨y, _⟩ => g y x
it3 fun f g x => match f x with | ⟨y, z⟩ => (match g y x with | ⟨w, _⟩ => ⟨w, z⟩)

```

The first candidate threads the state correctly. Now look at the second: it is type-correct and runs `g` on the *initial* state, throwing away whatever `f` did to it. That is the classic state-threading bug, and the type admits it just as happily. Types alone cannot tell these apart, which is why all candidates are shown and each is one keystroke from a test run. Continuation-passing style works too; this is the `bind` of the continuation monad:

session

```

λ> :synth ((A → R) → R) → (A → (B → R) → R) → (B → R) → R
it1 fun f g h => f (fun x => g x h)

```

With the inductive expansion of §5.5 the same game extends to data. Here is `Option.bind`, with the lazy `.none` candidate ranked first and the real one second. When in doubt, run one:

session

```

λ> :synth (∀ a b : Type, Option a → (a → Option b) → Option b)
it1 fun _ _ _ => .none
it2 fun _ _ x f => match x with | .none => (.none) | .some y => f y

```

5.3 Rank-N and impredicative goals

A polymorphic hypothesis is not just cargo: Djex instantiates it at types the goal itself supplies — including, under a guard, at *polymorphic* ones. Here the first candidate applies the identity hypothesis to the whole goal $Q \rightarrow Q$, an impredicative instantiation:

session

```

λ> :synth ((∀ p : Prop, p → p) → Q → Q)
it1 fun f => f _
it2 fun _ x => x

```

And a Church-encoded pair converts into a real conjunction — the quantified hypothesis is instantiated once at `p` and once at `q`, fed the matching projection each time:

session

```

λ> :synth (∀ p q : Prop, (∀ r : Prop, (p → q → r) → r) → p ∧ q)
it1 fun _ _ f => (f _ (fun x _ => x), f _ (fun _ y => y))

```

(Trailing candidates are elided.) Explicit $\forall$ binders — leading, nested, trailing, or interleaved — are woven into the candidate’s lambda automatically, and uses of quantified hypotheses get placeholder type arguments wherever Lean needs them (`f _ x`), so bounded rank-N candidates verify. Full impredicative inhabitation is undecidable; beyond the guard the answer is “no term found within bounds” and nothing stronger.

Djinn deliberately keeps three specialization families separate. The historical local family uses goal variables, opened skolems, premise scopes, and guarded quantified shapes. A final query-closed tail revisits only context-free schemes embedded in the requested goal and additionally admits closed, forall-free subtrees from that checked query. Loaded environment schemes remain a distinct family with retained source identity and may also use closed subtrees from loaded signatures.

The separation makes this provider-free goal reachable without treating the abstract `Mono` constant as a type variable:

session

```
λ> axiom QueryClosed.Mono : Type
λ> axiom QueryClosed.Token : Type
λ> axiom QueryClosed.Indexed : Type → Type
λ> :synth ((∀ a : Type, (a → QueryClosed.Token) → a →
...>   QueryClosed.Indexed a) → (QueryClosed.Mono → QueryClosed.Token) →
...>   QueryClosed.Mono → QueryClosed.Indexed QueryClosed.Mono)
  it1 fun f => f _
  it2 fun f g x => f _ (fun _ => g x) x
```

Both Djinn candidates instantiate `f` at `QueryClosed.Mono`; the second supplies a different valid constant callback. Standalone Exference returns the compact `fun f => f _`, and combined mode retains both Djinn spellings. The `synth-query-closed-rankn` transcript disables live-library premises, bounds Exference at 128 steps, and checks all three modes through final Lean 4.31 elaboration. A pure all-engine boundary regression covers the same projection, and remains in the suite alongside the later widening.

The new structural and nominal plans come after every established family, so historical result prefixes do not move. They carry historical local, loaded-scheme, and checked provider premises for composition, but do not reclassify loaded declarations as local schemes. The tail independently retains the six-binder, 16-per-scheme, 64-per-family, and 512-attempt caps. It is positive-only: exhaustion is `NoEvidence`, never a refutation. The detailed boundary is recorded in the `query-local closed-monotype` report.

Context-free instantiation chains extend to six leading binders. Leant inserts the inferred type arguments which Djinn leaves implicit, then asks Lean to verify every rendered candidate against the original type:

session

```
λ> axiom SixBinder.Six : Type → Type → Type → Type → Type → Type → Type
λ> :synth (∀ A B C D E F : Type, (∀ a b c d e f : Type, SixBinder.Six a b c d e f) →
...> SixBinder.Six F E D C B A)
  it1 fun _ _ _ _ _ x => x _ _ _ _ _
```

The focused `synth-six-binder-rankn` transcript disables live-library premises and pins that non-lexical source-order candidate through Lean 4.31 elaboration in standalone Djinn, standalone Exference, and combined mode. Its first query also discovers an active Lean instance and retains all six named quantified provider applications in every mode. Pure boundary tests separately retain the six-argument function-elimination shape and exact provider evidence across the same three engines. The five-binder transcript remains the

historical predecessor. The shared widening and its seven-binder controls are recorded in the [six-binder integration report](#).

The similarly named live synth-quartic-rankn regression has a subtler bridge shape than its Lean surface syntax suggests. After the four outer `FAll` binders for `Q`, `R`, `Z`, and `M`, each vacuous `forall A B C D E : Type`, `Q`-shaped component is serialized as five ordinary arrows from the shared opaque `Type` atom to its codomain, not as another `FAll`. Among the eight result leaves, only the four polymorphic identity leaves remain quantified. The regression therefore does not exercise the hypothesis-instantiation guard. Djex f3dd2495 lets standalone Exference forward those complete scoped arrow schemes through a structurally materialized nested product. The quartic submodule revision, c0c1a461, follows it with tuple-goal provenance: an independently scheduled structural route may use the eager whole-tree shortcut once, while fields emitted by the shallow alternative remain shallow. Equivalent recursive structural histories are suppressed without losing scoped or environment product reuse at depth. Exference admits the direct, independently checked candidate at search step 30. The unchanged 4096-step/1024-queue live run still ends with a bounded-tail note recording 39,308 queue prunes (the exact figure tracks the regenerated synth-quartic-rankn golden): search continued after the successful candidate and did not exhaust its ranked tail.

The live synth-quintic-rankn regression supplies the missing exact bridge case. It defines `Wide = forall A B C D E F G : Type, A × B × C × D × E × F × G`, so every binder is non-vacuous and remains an adjacent `FAll`. Leant 378f866 coalesces each uninterrupted `FAll` spine to one Djex `ForallType` binder list without crossing an `FInst`; the original fragment still supplies every explicitness slot to the renderer. One `Wide` therefore contributes one Djinn occurrence site rather than seven nested sites. With Djex d728719f, the first goal selects five non-contiguous opaque `Wide` results beside five opened identities at ten sites. The second selects five separated open identities beside six opaque `Wide` results at eleven sites, exercising the distinct dual plan. Leant 80f123a records both direct candidates after Lean 4.31 accepts them; neither live search is truncated. The sentinel now has seven binders, preserving the same occurrence coverage witness beyond the independent six-binder instantiation cap.

Vacuous local providers can retain the same query-supplied choice even when their result is an opaque Lean type. Here `Demo.Token` becomes a fixed ambient query rigid, not an unresolved inference variable:

session

```
λ> axiom Demo.Token : Type
λ> :synth ((∀ x : Type, x → x) → ({a : Type 1} → Demo.Token) → Demo.Token)
  it1 fun _ x => @x (∀ (a0_0 : Type _), a0_0 → a0_0)
λ> :set synth-engine exference
synth engine: exference
λ> :synth ((∀ x : Type, x → x) → ({a : Type 1} → Demo.Token) → Demo.Token)
  it1 fun f x => f _ (@x (∀ (a0_0 : Type _), a0_0 → a0_0))
```

This visible specialization is available only for a fully vacuous, context-free leading provider chain, and it selects only a closed proper type already present in the checked query. Fixed ambient rigids may occur in the provider result; a free flexible variable rejects the route. A local value has no global binder name, so Leant renders a positional `@` application. For a quantified argument it tries binder domains `_`, `Type _`, and `Prop`, in that order, and backend verification keeps the first spelling which elaborates. Named globals retain source-directed

applications. Each domain lane retains at most twelve site-and-style spellings, for a hard 36-variant bound only when the local quantified argument is domain-sensitive; ordinary duplicate spellings collapse to the historical smaller group.

An active Lean instance head can establish a provider-local choice even when the polytype does not occur in the query. In this session the goal contains only `Gap.Token`; the instance for `Gap.C` determines the quantified argument of the exact provider:

```
session
λ> axiom Gap.Token : Type
λ> class Gap.C (a : Type 1) : Prop where witness : True
λ> instance : Gap.C (∀ x : Type, x → x) := (True.intro)
λ> axiom Gap.polyGlobal {a : Type 1} [Gap.C a] : Gap.Token
λ> :set synth-engine djinn
synth engine: djinn
λ> :synth Gap.Token
  it1 Gap.polyGlobal («a» := (∀ (a0_0 : Type _), a0_0 → a0_0))
λ> :set synth-engine exference
synth engine: exference
λ> :synth Gap.Token
  it1 Gap.polyGlobal («a» := (∀ (a0_0 : Type _), a0_0 → a0_0))
λ> :set synth-engine both
synth engine: both
λ> :synth Gap.Token
  it1 Gap.polyGlobal («a» := (∀ (a0_0 : Type _), a0_0 → a0_0))
```

An active head can also determine a higher-kinded binder which is absent from the provider body. Leant preserves the argument's bounded `Type → ... → Type` kind, so both checked engines keep this constraint-only application visible:

```
session
λ> namespace Higher
λ> axiom Wrap : Type → Type
λ> class VacuousChoice (F : Type → Type) : Prop where witness : True
λ> instance : VacuousChoice Wrap := (True.intro)
λ> axiom VacuousToken : Type
λ> axiom vacuous {F : Type → Type} [VacuousChoice F] : VacuousToken
λ> end Higher
λ> :set synth-engine djinn
synth engine: djinn
λ> :synth Higher.VacuousToken
  it1 Higher.vacuous («F» := Higher.Wrap)
λ> :set synth-engine exference
synth engine: exference
λ> :synth Higher.VacuousToken
  it1 Higher.vacuous («F» := Higher.Wrap)
```

One exact provider may retain several distinct successful instance-head vectors. In the same transcript, heads for `AlternativeChoice Wrap` and `AlternativeChoice (Pair Nat)` produce exactly `Higher.alternative («F» := Higher.Wrap)` and `Higher.alternative («F» := (@Higher.Pair Nat))`. Standalone Djinn ranks `Wrap` first, standalone Exference ranks `Pair Nat` first, and combined mode uses Djinn's order after stable exact-spelling dedu-

plication. That is an engine-ranking difference: every mode must retain both exact alternatives once. The transcript also carries one heterogeneous two-binder vector whose positional kind arities are one and two, and requires `Higher.multiVacuous («F» := Higher.Wrap) («G» := (@Higher.Triple Nat))` in all three modes.

Discovery opens no more than six leading type binders on one exact provider while retaining its erased instance constraints. It inspects at most 32 active heads per provider in Lean resolver order. Each selected head and its complete constraint closure run under restored metavariable state. The seed head stays fixed while its returned instance subgoals and every other provider constraint are solved in that same context; a failure or unresolved binder rejects the head as a whole. One success contributes one complete vector in leading-binder order, never an uncorrelated scalar pool. For each argument, discovery also reduces its type to the admitted first-order kind language and retains the number of Type arrow domains. Zero denotes proper type; positive counts preserve bare and partially applied constructors.

Leant retains at most 16 alpha-distinct vectors per provider and passes no more than 32 provider/vector associations to Djex. Each vector has exact provider arity and at most six arguments. The command takes that aggregate prefix before an argument affects family planning, rigidity, or translation. Live provider metadata uses `(instantiations (args (kinded N ...) ...))`. For context-free sources, Leant folds each bounded arrow count into a Djex `GroundKind`, pairs it with the translated type, and calls the checked kinded Djinn or Exference assignment adapter. An exact contextual Exference vector whose complete fact group survives the replayed inventory trial instead contributes ground facts and does not also enter that context-free adapter. If no group survives for an exact-evidence provider, its successfully translatable bounded vectors take the historical context-erased fallback. Djinn always uses that erased compatibility projection.

Live discovery, wire parsing, and engine filtering all admit at most 64 Type-arrow domains, whose simple right-associated kind has $2 * 64 + 1 = 129$ constructors. After a context-free assignment passes Djex's provider, scheme, context, and exact-arity checks, the pinned adapters productively preflight all of that assignment's supplied kinds before recursive kind inference, same-provider comparison, kind conversion, or paired type elaboration. Oversized and cyclic caller-built kinds therefore fail finitely. Metadata-free and binder-only inventories remain valid. Historical exact arguments stay in their original vectors and default each position to proper kind; the `(candidates ...)` form parses as legacy proper-kind unary vectors and can never authorize a ground class fact.

A vector containing depth truncation (`FDepth`) or a legacy raw instance marker (`FInst`) is discarded in full. Supported provider-source and exact-assignment serialization may instead retain a bounded `FExactContext`: the exact nominal Lean class head, every ordered argument's bounded first-order kind arity, and the contextual body. A positive-arity source argument may retain its enclosing `FA11` variable bare or partially applied only to proper-type arguments. A ground assignment instead retains a canonical bare or partially applied nominal head with proper-type supplied arguments, serialized as `(kinded N (nominal "Head" ...))`. Historical fragment arguments remain readable for compatible nonstructural heads, but carry no structural identity. Canonical nominal `Prod` and `Sum` heads are supported only at total arity two: direct occurrences use Djex's boxed-pair and `Either` identities, and contextual occurrences preserve that same identity. Legacy structural payloads and the other six structural heads remain fail-closed. Malformed, open, dictionary-dependent, term-indexed,

and otherwise unsupported contexts still discard their complete vector. A bounded pre-scan also discards only assignment vectors that touch inconsistent class-kind or nominal-family arity claims. Provider-prefix fallback keeps metadata attached to its exact declaration. The checked runners validate provider identity, arity, supplied positional kinds, closure, and context before consuming each vector once without Cartesian reconstruction. Thus a later or alpha-identically typed provider cannot donate its assignment to an earlier one.

Leant also retains applications whose arguments are proper types. A bound first-order constructor remains a higher-kinded variable in Djex; an opaque Lean family receives one rigid nominal head across its occurrences. Both engines can therefore see a query-supplied polytype without learning any hidden implementation:

session

```
λ> axiom Wrap : Type 1 → Type
λ> :synth ((∀ a : Type 1, Wrap a) → Wrap (∀ b : Type, b → b))
  it1 fun x => x _
λ> :set synth-engine exference
synth engine: exference
λ> :synth (∀ (F : Type 1 → Type), (∀ a : Type 1, F a) → F (∀ b : Type, b → b))
  it1 fun _ x => x _
```

The serializer admits only first-order kind arrows between universes and only application arguments whose own type is a universe. Term-indexed families such as `P 3` remain opaque. Canonical nominal `Prod` and `Sum` are admitted as bare or partially applied higher-kinded assignment heads only when their supplied-count plus residual kind arity is exactly two. Direct assignments use the same boxed-pair and `Either` identities as saturated structural translation. Legacy structural payloads cannot claim that identity; `And`, `PProd`, `Or`, `PSum`, `Iff`, and `Not` remain excluded as higher-kinded assignment heads. Saturated uses keep their structural translation. A retained nominal application can support a positive, Lean-verified candidate but always poisons a negative verdict because its constant hides structure. Query-derived candidate types come only from proven proper-type positions below arrows and tuples in the sequent, including quantified subtrees already present there; provider-local assignment arguments may also come from the bounded active-instance channel above. This is guarded, evidence-directed impredicativity, not general second-order search. The query-derived visible-provider route is capped at six binders and 32 argument combinations; Djex's internal Haskell-instance route remains monotype-only. The external Lean evidence route has the separate six-binder, 32-head, 16-vector-per-provider, and 32-vector limits. Open, incomplete, wrong-arity, depth-truncated, malformed, and unsupported vectors fail closed. Closed quantified arguments may contain a structured exact-assignment-local nominal class context; raw `FInst`, free or term-indexed contexts, dictionary-dependent result bodies, and chains with seven or more leading binders remain deliberate incomplete boundaries. The original local-provider contract, scalar active-instance extension, and correlated follow-up are recorded in the `scoped local-provider report` and the `provider-local instance-head report`, followed by the `correlated assignment report`, with the shared successor recorded in the `five-binder integration report`, the current boundary in the `six-binder integration report`, and the additive local tail in the `query-local closed-monotype report`, and the structured context

boundary in the contextual provider-assignment report, and the structural higher-kinded extension in the higher-kinded contextual-assignment report.

Unit regressions pin the kinded wire format, kind and argument order, finite bounds, whole-vector deduplication, and the exact vacuous applications under Djinn, Exference, and combined mode. The live `synth-provider-higher-kind-assignment` transcript retains the mixed higher-kinded/rank-N vector, the heterogeneous arity-one/arity-two multi-vacuous vector, and both exact `Wrap` and `Pair Nat` alternatives for one provider. Engine-specific ordering remains ranking policy. The live `synth-provider-structural-assignment` transcript additionally pins direct `Prod`/partial-`Sum` specialization and a closed contextual assignment containing `Binary Prod` and `Unary (Sum Fixed)` in Djinn, Exference, and combined mode.

5.4 Impossibility, proved

When Djinn’s complete search exhausts a fully translated goal, failure is a theorem:

session

```
λ> :synth (∀ a b : Type, a → b)
provably uninhabited – no closed term of this polymorphic type exists
```

No Lean tactic gives you this answer (a failing `exact?` proves nothing). Note the careful wording: the verdict is about *closed terms of the polymorphic type*. It does not say any particular instance $a \rightarrow b$ is empty (choose $a = b$ and it isn’t), and for propositions it is about *constructive* provability, a distinction §5.7 exploits. When the translation had to hide structure behind an opaque atom, the verdict backs off to match:

session

```
λ> :synth (∀ a : Type, Sigma (fun _ : Nat => a))
no term found within bounds (opaque atoms `Nat` hide structure the engine cannot
↳ analyze – not a refutation)
```

5.5 Inductive types

A non-recursive, non-indexed inductive or structure applied to all its parameters is expanded into a generalized sum of products: constructors become introduction rules, case analysis the elimination rule. This covers the standard library’s little workhorses —

session

```
λ> :synth Ordering
  it1 .lt
  it2 .eq
  it3 .gt
λ> :synth (∀ a : Type, Option (Option a) → Option a)
  it1 fun _ _ => .none
  it2 fun _ x => match x with | .none => (.none) | .some y => y
λ> :synth (∀ e a b : Type, Except e a → (a → b) → Except e b)
  it1 fun _ _ _ x f => match x with | .error y => .error y | .ok z => .ok (f z)
```

— `Option.join` and `Except.map` synthesized rather than remembered. It extends to Type-valued classes-as-data like `Decidable`, where the eliminations write themselves, up to and including the instance combinators a library author would otherwise write by hand:

session

```
λ> :synth (∀ p : Prop, Decidable p → Decidable (¬ p))
  it1 fun _ x => match x with | .isFalse k => .isTrue k | .isTrue y => .isFalse (fun
    ↪ k1 => k1 y)
λ> :synth (∀ p q : Prop, Decidable p → Decidable q → Decidable (p ∧ q))
  it1 fun _ _ x y => match x with | .isFalse k => .isFalse (fun {z, _} => k z) |
    ↪ .isTrue w => (match y with | .isFalse k1 => .isFalse (fun {_, x1} => k1 x1) |
    ↪ .isTrue x2 => .isTrue {w, x2})
λ> :synth (∀ p q : Prop, Decidable p → Decidable q → Decidable (p → q))
  it1 fun _ _ x y => match x with | .isFalse k => .isTrue (fun z => absurd z k) |
    ↪ .isTrue w => (match y with | .isFalse k1 => .isFalse (fun f => k1 (f w)) |
    ↪ .isTrue x1 => .isTrue (fun _ => x1))
```

The last two are decidability of conjunction and of implication, assembled by case analysis on both instance arguments: four branches each, every branch carrying its own little proof.

Session-declared types participate the moment you declare them, and refutations over expanded inductives remain sound, because the engine saw the complete constructor list:

session

```
λ> structure Pair (A B : Type) where
...>   fst : A
...>   snd : B
...>
λ> :synth (∀ a b : Type, a → b → Pair a b)
  it1 fun _ _ x y => {x, y}
λ> :synth (∀ a b : Type, Pair a b → Empty)
provably uninhabited – no closed term of this polymorphic type exists
```

Recursive types cannot be expanded; elimination is where undecidability lives. Their *constructors* are still sound introduction rules, though, so building is possible even where analysis is not:

session

```
λ> :synth Nat
  it1 .zero
λ> :synth (Nat → Nat)
  it1 fun x => x
  it2 fun x => Nat.add x x
  it3 .succ
  it4 fun _ => .zero
λ> :synth (∀ a : Type, a → List a)
  it1 fun _ _ => .nil
```

5.6 Recursion from the library

:synth will never invent a `Nat.rec`-based program, but for the everyday recursive types it does not have to: the library already wrote the recursion. A goal that mentions `List` or `Nat` brings a rated inventory of library functions with it — `List.map`, `List.foldr`, `List.append`, `List.flatten`, `List.length`, `List.replicate`, `Nat.add`, ... — instantiated at the goal's own types and offered to the engine as extra premises. The enumeration prefers proofs that use the goal's own arguments, generally putting a direct library answer first while retaining distinct choices between same-typed arguments:

session

```
λ> :synth ((a → b) → List a → List b)
  it1 fun f x => List.map f x
  it2 fun f x => List.reverse (List.map f x)
λ> :synth (List (List a) → List a)
  it1 fun x => List.flatten x
  it2 fun x => List.reverse (List.flatten x)
λ> :synth (Nat → a → List a)
  it1 fun x y => List.replicate x y
  it2 fun x y => List.reverse (List.replicate x y)
λ> :synth (List a → List b → List (a × b))
  it1 fun x y => List.zip x y
  it2 fun _ _ => .nil
λ> :synth ((a → b → c) → List a → List b → List c)
  it1 fun f x y => List.zipWith f x y
  it2 fun _ _ => .nil
λ> :synth (List a → List a → List a)
  it1 fun x _ => x
  it2 fun _ x => x
  it3 fun x _ => List.reverse x
  it4 fun x y => List.append x y
  it5 fun x y => List.append y x
```

(Trailing candidates and the truncation notes are elided here.) The inventory is a ratings list in Djex's `*.ratings` format — lower is better, 100 or more disables — and a project file `leant.ratings` (lines of `Name Rating, # comments`) merges over the defaults at startup, so re-ranking, disabling, or growing the inventory is editing a list, not writing code.

The library search runs beside the constructor search of the previous section, in a mode where the recursive occurrences are sealed atoms: without `nil/cons` in play, every candidate must route through the goal's own arguments and the offered functions, which is what keeps `List.map` ahead of the endless supply of closed terms that ignore the input. Both searches' survivors are shown together, library candidates first; `:set synth-library off` restores the constructors-only behavior. As always, the premises are only *offers*: every candidate is re-elaborated by the backend, so a useless premise costs search time, never a wrong answer, and negative verdicts never come from the library run at all.

5.7 Classical candidates

Constructive logic goes further than newcomers expect, and :synth stays inside it whenever it can. Triple negation collapses, and contraposition holds in one direction, with no axioms in sight:

session

```
λ> :synth (∀ p : Prop, ¬¬¬p → ¬p)
  it1 fun _ k x => k (fun k1 => k1 x)
λ> :synth (∀ p q : Prop, (p → q) → ¬q → ¬p)
  it1 fun _ _ f k x => k (f x)
```

Pierce’s law is different: it has no constructive inhabitant, and :synth can prove that. But rather than stopping at the refutation, it first gives constructive live providers a chance to supply a Lean-verified term. Only if those lanes fail does it ask the classical question, and answer with a term whose spelling shows what was used:

session

```
λ> :synth (∀ p q : Prop, ((p → q) → p) → p)
  it1 fun _ _ f => match Classical.em _ with | .inl x => x | .inr k => f (fun y =>
    ↳ absurd y k)
λ> :synth (∀ p : Prop, ¬¬p → p)
  it1 fun _ k => match Classical.em _ with | .inl x => x | .inr k1 => absurd k1 k
```

These are the proofs a human would write: one case split on excluded middle, one absurd where the contradiction lands. The hard direction of De Morgan needs excluded middle twice, once per disjunct, and the candidate reads as exactly that case analysis:

session

```
λ> :synth (∀ p q : Prop, ¬(¬p ∧ ¬q) → p ∨ q)
  it1 fun _ _ k => match Classical.em _ with | .inl x => .inl x | .inr k1 => (match
    ↳ Classical.em _ with | .inl y => .inr y | .inr k2 => absurd {k1, k2} k)
```

Less famous gaps between the two logics come out the same way. Weak excluded middle, and the Gödel–Dummett formula:

session

```
λ> :synth (∀ p : Prop, ¬p ∨ ¬¬p)
  it1 fun _ => match Classical.em _ with | .inl x => .inr (fun k => k x) | .inr k1 =>
    ↳ .inl k1
λ> :synth (∀ p q : Prop, (p → q) ∨ (q → p))
  it1 fun _ _ => match Classical.em _ with | .inl _ => (match Classical.em _ with |
    ↳ .inl x => .inl (fun _ => x) | .inr k => .inr (fun y => absurd y k)) | .inr k1 =>
    ↳ .inl (fun z => absurd z k1)
```

Two searches stand behind this. First, one excluded-middle assumption per atomic subformula is handed to the intuitionistic engine — complete for the propositional fragment, since intuitionistic logic plus atom instances of $p \vee \neg p$ is classical propositional logic. The findings are rendered as `Classical.em` case splits. If that misses, the goal’s double-negation

translation runs instead (complete by Glivenko’s theorem) and the candidate is wrapped in `Classical.byContradiction`. Both kinds are verified like any other candidate. Excluded middle remains one six-group batch. An explicit behavioral filter uses the same command-local bank in both routes: either no verification or complete all-rejection enters double negation, whose ordinary cursor policy may inspect 12+12 standalone or 24+24 combined groups from one engine outcome. Any survivor or preserve-all result stops, with no third cursor probe. Filtering retains the original command deadline through both routes. Rank and disabled commands instead preserve their historical classical budgets by capturing a fresh configured duration separately at each route they actually reach. The sound-refutation caller accumulates every completed batch and finalizes only once. The whole behavior sits behind a switch:

session

```
λ> :set synth-classical off
synth classical: off
λ> :synth (∀ p : Prop, p → ¬ p)
provably uninhabited – no closed term of this polymorphic type exists
(constructively – a classical proof may still exist; this is not a disproof of the
  ↪ proposition)
```

5.8 Dependent formulas as cargo

Leant’s translation makes no attempt to analyze dependent types. It does not refuse them either: a dependent subformula becomes an opaque *atom*, compared up to α -equivalence, which the engine can carry from one side of a goal to the other. Transport works; analysis does not:

session

```
λ> opaque P : Nat → Prop
λ> opaque Q : Prop
λ> :synth ((∀ n : Nat, P n) → Q → Q → (∀ n : Nat, P n))
  it1 fun (x, y) => {y, x}
```

The engine never looked inside $\forall n, P n$; it swapped a sealed box. A goal that would require opening the box (an induction, a rewrite, a case split on an index) is refused with a reason; that work belongs to `:prove`.

5.9 Synthesis inside a proof

Bare `:synth` in prove mode targets the current goal *with its hypotheses as premises*. Candidates are offered as `it1`, `it2`, ... exactly as in the session, except that here `itN` splices the candidate applied to your hypotheses, so exact `it1` closes the goal. The hypotheses can carry real weight; below, the case analysis on the `Decidable` instance comes from `hd`:

session

```
λ> :prove ∀ p q : Prop, Decidable p → (p → q) → (¬p → q) → q
entering prove mode – type tactics; :help for commands
```

```

┬- ∀ (p q : Prop) (a : Decidable p), (p → q) → (¬p → q) → q
suggestion: exact fun p q a a_1 a_2 => dite p a_1 a_2 (closes the goal)
┬-> intro p q hd hpq hnq
p q : Prop
hd : Decidable p
hpq : p → q
hnq : ¬p → q
┬- q
suggestion: exact dite p hpq hnq (closes the goal)
┬-> :synth
(synthesizing with hypotheses p q hd hpq hnq as premises)
  it1 fun _ _ x f g => match x with | .isFalse k => g k | .isTrue y => f y
┬-> exact it1
All goals accomplished
finish with :qed [NAME], inspect with :script
┬-> :qed byCases
saved: theorem byCases : ∀ p q : Prop, Decidable p → (p → q) → (¬p → q) → q

```

Note the contrast with the tactic suggestion above it: the suggestion *found* the library's `dite`, while `:synth` *composed* the case split from the goal's own material. Both are verified; they simply answer different questions.

It also slots into an ordinary multi-goal proof. Split a conjunction with `constructor` and `let :synth` deal with each half in turn. Notice the prompt tracking the number of open goals, and the projections from `h` showing up without being asked for:

session

```

λ> :prove ∀ p q r : Prop, (p → q) ∧ (q → r) → p → r ∧ (p → q)
entering prove mode – type tactics; :help for commands
┬- ∀ (p q r : Prop), (p → q) ∧ (q → r) → p → r ∧ (p → q)
suggestion: intro p q r h h1 <=> simp_all (closes the goal)
┬-> intro p q r h hp
p q r : Prop
h : (p → q) ∧ (q → r)
hp : p
┬- r ∧ (p → q)
suggestion: simp_all only [forall_const, imp_self, and_self] (closes the goal)
┬-> constructor
– goal 1 of 2 –
case left
p q r : Prop
h : (p → q) ∧ (q → r)
hp : p
┬- r
– goal 2 of 2 –
case right
p q r : Prop
h : (p → q) ∧ (q → r)
hp : p
┬- p → q
suggestion: simp_all only [forall_const] (closes the goal)
┬-2> :synth
(synthesizing with hypotheses p q r h hp as premises)
  it1 fun _ _ _ (f, g) x => g (f x)

```

```

|-2> exact it1
case right
p q r : Prop
h : (p → q) ∧ (q → r)
hp : p
|- p → q
suggestion: simp_all only [forall_const, imp_self] (closes the goal)
|-> :synth
(synthesizing with hypotheses p q r h hp as premises)
  it1 fun _ _ => {f, _} x _ => f x
|-> exact it1
All goals accomplished
finish with :qed [NAME], inspect with :script
|-> :qed chain
saved: theorem chain : ∀ p q r : Prop, (p → q) ∧ (q → r) → p → r ∧ (p → q)

```

Unlike `exact?`, which finds an *existing* lemma, this composes a new term from the goal's own material: a constructive complement to the finisher tactics, needing no premise database and no imports. The division of labor that emerges is pleasant: the genuinely dependent moves (quantifiers, induction, rewriting) stay yours, and the propositional bookkeeping between them does not.

The classical fallback of §5.7 follows `:synth` into prove mode. Double-negation elimination has no constructive proof, so the candidate arrives as an excluded-middle case split, and `:qed` turns it into a named theorem — no tactic beyond `exact` was typed:

session

```

λ> :prove ∀ p : Prop, ¬¬p → p
entering prove mode – type tactics; :help for commands
|- ∀ (p : Prop), ¬¬p → p
suggestion: exact fun p a => Classical.byContradiction a (closes the goal)
|-> :synth
  it1 fun _ k => match Classical.em _ with | .inl x => x | .inr k1 => absurd k1 k
|-> exact it1
All goals accomplished
finish with :qed [NAME], inspect with :script
|-> :qed not_not_elim
saved: theorem not_not_elim : ∀ p : Prop, ¬¬p → p

```

5.10 Engines and budgets

The default engine is Djinn: complete, terminating, and the only source of refutations. `:set synth-engine exference` switches to the ranked heuristic search, and both runs the two together — negative verdicts still come only from Djinn. One standalone batch sends at most 12 fresh candidate groups to Lean. A combined batch gets 24 and retains both standalone frontiers: writing D and E for fresh Djinn and Exference groups, its order is D1–D4, E1–E12, D5–D12, then alternating tails. Within each Exference invocation, Leant stable-deduplicates rendered groups before applying the internal 60-candidate collection window. The first

spelling remains authoritative, while repeated backend histories cannot hide later distinct terms behind the bound; the outer 12- and 24-group batch frontiers then apply.

`Leant.Synth.Engine` exports an opaque `DetailedSynthCursor` over one lazy detailed outcome, and `Main` now consumes it for bounded filter progression. For a nonempty candidate outcome below the cap, an admitted request of 1 through 60 returns a nonempty ordered batch and a lazy successor; cumulative observation stops at the existing product-wide 60-group hard cap. Natural exhaustion is distinct from reaching that cap: after group 60 the cursor reports the cap without probing group 61, even when the finite stream ended exactly there. Its force helper traverses the selected group-list and variant-list spines and their constructors far enough to project text, forces each selected route to weak head normal form, traverses each selected rendered spelling's `String` list spine through `length`, and traverses the run-note list spine and each note's `String` spine through `length`. It neither forces `Char` values nor performs text encoding, and leaves semantic sidecars, the successor, and the unselected tail lazy for a caller-owned deadline. `Main` imports the opaque cursor type without its constructor, the visible step taxonomy, batch accessors, and cursor operations. It neither names the hidden batch type nor imports either hidden constructor.

Ordinary ranking retains its historical verifier stop after five accepted groups and consumes one batch. An explicit behavioral filter verifies and assesses one complete 12-group standalone or 24-group combined batch. Only no verification or complete behavioral rejection may consume one successor of the same width from that cursor, giving 12+12 or 24+24 groups from one engine run. Every batch is verified and assessed exactly once through the same context; `Main` neither reruns the engine nor reassesses the first batch. `Survivor` and `preserve-all` stop immediately, and a continuing second batch stops without a third tail probe. In either operation at most five candidates are displayed and bound. `Exference` runs under explicit budgets (`:set synth-steps`, default 4096), and truncation is always reported:

session

```
λ> :set synth-engine both
synth engine: both
λ> :synth ((A → B → C) → (A → B) → A → C)
  it1 fun f g x => f x (g x)
note: exference: search truncated: step limit reached, queue limit pruned 36194
```

Every engine mode first gives a structurally accepted goal a provider-free lane and asks Lean to verify its candidates. When no provider-free term verifies, or an authorized behavioral filter rejects every verified occurrence, Lean discovers at most 80 live term providers under the namespaces used by the target (plus exact session declarations). The command's filter bank and accumulated rejection rows cross those lanes. A complete Djinn refutation is retained as a sound control-flow fallback while constructive provider lanes run: the first survivor or `preserve-all` result wins; an empty, all-rejected, unavailable, timed-out, or unsuccessful provider search restores the refutation before any classical retry, without discarding completed lane accounting. The focused `synth-provider-refutation-fallback` transcript checks exact rank-N overrides under Djinn, `Exference`, and combined mode. It also enables classical search for a Peirce-shaped goal and confirms that an exact constructive provider wins first, while a no-provider control retains the original sound refutation. Pure tests pin the same direct, rank-N, and unusable-provider boundaries below the REPL layer. Type constructors

and type families, whose fully peeled result is a sort, are excluded before search. Exact-result session and public declarations lead the ranking; conventional implementation-worker names remain eligible but follow public fallbacks, while session declarations bypass that naming heuristic. Exference searches the ranked inventory directly; Djinn tries discovery-order prefixes of 1, 4, and 16 providers before the full bounded inventory, and may specialize a loaded scheme at a closed monotype or guarded rank-N polytype. In combined mode both engines run the singleton and full lanes, while the intermediate prefixes are Djinn-only. The serializer derives a canonical query from the sorted, deduplicated namespace roots and final result head, rather than from the goal's raw spelling. Successful inventories — including an empty one — are kept in a generation-aware 12-entry LRU. Imports, ordinary declarations, and session-restoring operations invalidate that generation; generated `it1`, `it2`, ... bindings cannot be providers and preserve it.

For an exact provider with erased instance constraints, the same discovery pass can attach bounded active-instance-head assignments. Resolution follows Lean's own instance order and isolates the selected head plus its complete constraint closure. One successful head yields one ordered vector; an incomplete or unusable head yields nothing. Provider slicing retains the vector only with its source declaration; after whole-vector context/depth and arity filtering Leant takes the command-wide prefix before planning or translation. Each argument crosses with its bounded Type-arrow kind; Leant reconstructs a Djex GroundKind and calls the checked kinded Djinn or Exference assignment runner. Leant rejects residual kind arities above 64 before that bridge, and pinned Djex independently rejects a supplied GroundKind above 129 constructor nodes before recursive operations on that assignment. Lean still reconstructs every dictionary during final elaboration—only the closed kind/type vector crosses the bridge.

Exact rendered spellings keep their first scheduled occurrence. Before a later provider lane is forced or capped, every spelling in an earlier complete no-verified or all-behaviorally-rejected bounded frontier is removed from each engine's source stream and newly empty groups are dropped. Rediscovered candidates therefore spend none of that lane's 12- or 24-group fresh frontier.

An entirely optional last stage can consult Z3. Passing `--length-ranking-config` at startup (§2) activates finite-list-spine *Length* behavioral assessment: a passive JSON contract states a law over the lengths of list-spine arguments and results, eligible verified candidates are translated into canonical QF_LIA queries, and at most one scoped Z3 worker is opened per batch when pure replay does not finish the assessment. Only a counterexample that Leant's vendored Djex engine has independently replayed against the exact checked problem can move a candidate. In the default rank operation such candidates are stably demoted, never pruned. An explicit `--behavior-mode filter` operation instead omits exactly those replayed-counterexample occurrences, prints each one separately as rejected, and binds only survivors. A raw `sat`, `unsat`, or `unknown` has no authority at all; neither do a preparation refusal, unassessed input, heuristic status, completed finite box, or established applicable domain.

Every assessment remains deliberately batch-bounded. A rank request asks Lean for at most five successful groups and consumes one engine batch. A filter request may verify and assess two 12-group standalone or two 24-group combined batches from one cursor. The second is admitted only after no verification or complete all-rejection. Excluded middle

stays one six-group batch; double negation uses the ordinary at-most-two-batch policy and a filter-mode Djinn candidate cutoff of 60. Main introduces one rank-2 context before goal translation and carries it through both batches, universe retries, and every ordinary, provider, and classical route. Each verified frontier enters exactly one assessment, while the filter context's one nominal scalar or product bank can replay an earlier counterexample against a later occurrence, successor batch, or same-scope lane. Existing ranking and direct Selection compatibility paths keep the raw four-vector MRU unchanged; rank contexts have no bank. Only after a terminal selection does Leant take at most five survivor presentations for display and binding; every accumulated rejection is still printed.

The exact option-bearing grammar is `:synth [--behavior-mode rank|filter] [--length-contract PATH] -- TYPE`; mode must precede contract and the standalone delimiter is mandatory. Ordinary `:synth TYPE` remains delimiter-free. Omitting behavior mode selects rank. A per-command contract swaps in one passive scalar-or-pair assertion for that command; filter without a contract uses the startup contract. Filtering and every contract form require the already activated startup policy and are rejected before contract-path admission or file IO when it is absent. The startup grammar is one current strict, versionless shape selected by `rankingDomain`; the separate contract-only grammar is likewise one current strict, versionless `{format, rankingDomain, contract}` shape. It is parsed once with no migration fallback and carries only the command-local passive scalar or binary-product contract, never startup execution policy. The startup policy fixes one guarded recursive piecewise-affine applicable-domain algorithm. Programmatic callers have one short builder, `enableLengthRankingApplicableDomainValidation`, and the current assessments are `ApplicableDomainEstablished` and `LengthSpinePairApplicableDomainEstablished`. Presentation uses only `renderLengthApplicableDomainValidationNote` and `renderLengthSpinePairApplicableDomainValidationNote`. Djex keeps the direct, affine, relational, strict, quotient, extrema, monus, Boolean-union, and atomic analyses as private ordered fallbacks; they are not separately selectable Leant contracts. Its recursive expression grammar includes an all-or-nothing `LengthIf`: both condition polarities and both arms must be fully supported, the true arm precedes the false arm, and raw guard alternatives consume the existing branch budget before contradictions collapse. For example, `(if x <= 2 then x else 5) <= 3` establishes the provider-independent box `[[2]]` with three visits, three unique assignments, and three applicable assignments entirely before process IO; an unsupported arm makes the whole fallback atom inapplicable. The original checked problem is still replayed over every unique assignment and remains semantic authority. The former public aliases were deleted without migration because Leant is experimental and promises no stability or backward compatibility. This API reset does not alter the versionless file shapes or the seven numeric limits inside `applicableDomainValidation`. The selected-bank, domain, origin, deferred opening, live query, post-unsat input-box, scoped-budget, and stable preference boundaries remain unchanged. Ranking and direct Selection compatibility entrances keep the four-vector MRU local to one verified batch. The vendored Djex revision separately exposes candidate-independent scalar and binary-product bank scopes, bounded immutable input stores, and exact query-owned fresh replay. Leant's package-private adapter now includes nominal scalar/product contexts over that bridge: each retains validated limits and at most one active scope, resets on exact scope drift, threads authoritative charged successors through newest-first replay, and separates a replay hit from evaluation-free promotion. Context

transitions are serialized in exception-restoring cells; the successor state is fully forced and the selected outer classification is forced to weak head normal form before commit, so an unexpected exception leaves the old state and propagates.

Only the additive context-aware Configuration, Selection, and internal Ranking path uses those contexts, and Integration selects that path only for filtering. An unsimplified bank hit is promoted without reevaluation; a simplified hit is freshly recorded under simplification origin. Fresh live and solver-independent receipts retain their corresponding cache origins only without a strict reduction; every strict reduction records under simplification origin. The candidate keeps its pre-record final receipt. Completed cache transitions survive a later candidate-report fallback.

Integration introduces one reusable rank-2 LengthAssessmentContext; repeated context assessments reuse its filter bank, while rank contexts contain no bank. Main now opens one context inside synthRun before initial translation and supplies it to every cursor batch and later route. Its fresh-per-call compatibility request assessor remains for other callers. The rank-2 owner dies when the command returns and cannot enter REPL state.

Filtering runs the same scalar or nominal pair assessor, maps every report back to its exact callback occurrence by the retained original index, and seals one bounded total partition. Both survivors and rejections are restored to original callback order; filter mode does not inherit ranking order. Any batch-wide post-verification-adaptor, ranking, evidence-association, admission-limit, or partition-seal failure preserves the complete original verified batch and exposes no partial rejection; the warning path still displays at most five original candidates. That preserve-all result does not roll back bank replay, promotion, or recording transitions which already completed. A surviving bounded-positive receipt keeps its established note. A rejected occurrence keeps its exact counterexample or simplification note. If all verified occurrences are rejected, the batch is handled but not yet finalized: like a no-verified batch, it may consume the one permitted same-run successor or enter the next provider or classical route with the same bank. Survivor and preserve-all results remain terminal.

Main wraps each lazy engine outcome once in the opaque cursor. Rank and disabled commands admit one candidate batch. Filter-mode ordinary, provider, and double- negation runs admit a second same-width batch only after no verification or all-rejection: at most 12+12 groups for a standalone engine and 24+24 for both. Excluded middle remains one six-group batch. The second batch ends the run without a third tail probe. Every batch is verified and assessed once through the command context; the cursor successor never reruns an engine.

Each batch outcome keeps its complete bounded rendered-spelling frontier separate from the exact variants which reached Lean callbacks. One private run receipt restores the at-most-two outcomes to chronological order, concatenates their complete frontiers for provider deduplication, and preserves cumulative group counts and debug ordinals. It then prepends the outcomes to the existing lazy reverse command accumulation. Verification and assessment perform no candidate presentation or synthesis-splice mutation. A pure fold distinguishes no verified receipt, survivors, aggregate all-rejection, and preserve-all failure, stopping at the first terminal result.

Ordinary run notes attach once to the chronologically rightmost handled batch. If every batch is no-verification, they remain only in the run receipt for the final diagnostic. Classical handled outcomes remain note-free. One deferred finalizer owns metrics, warnings, binding

and cache effects, rows, and handled notes. Terminal candidate rows print first; rejection rows are every earlier all-rejected projection followed by the terminal batch's rejections. Final all-rejection clears the old splice cache once; survivor or preserve-all output replaces it once.

Continuation does not fill a five-survivor quota within a run or across lanes. The first batch with even one survivor stops scheduling, ranking retains its historical five-success quota, and filtering retains each 12/24 batch width and the single six-group EM width. On ordinary all-rejected exhaustion the accumulated rejection output suppresses an unrelated no-term or Lean-verification diagnostic. Without a retained sound refutation, timeout, cursor-admission failure, or engine error finalizes completed outcomes first and then prints its corresponding abnormal diagnostic.

Filter mode uses the original absolute command deadline for every constructive run, both batches, excluded middle, and double negation; it does not recapture a clock at a classical boundary. Rank and disabled modes keep that command deadline for ordinary and provider work but preserve their historical classical budgets: each reached excluded-middle and double-negation route independently captures a fresh configured duration. A skipped route captures nothing, and timeout zero remains unbounded.

A provider-free sound Djinn refutation still masks provider timeout/error and controls the later classical path, but completed provider outcomes stay in the accumulation. When the Glivenko split exists, excluded-middle no-verification or all-rejection enters double negation; survivor or preserve-all stops. The sound-refutation caller finalizes once. Aggregate no-verification keeps the existing provably-uninhabited claim, unresolved-universe annotation, and constructive/classical hint in that order. Aggregate all-rejection suppresses the unrelated claim and hint while keeping the goal-qualified annotation.

Behavior mode, assessment contexts, selection results, both kinds of replay-input store, accumulation, and lane outcomes are never persistent REPL state. This is a bounded command-local reuse and continuation slice, not quota filling or the proposal's complete CEGIS loop: no engine accepts a counterexample-directed replacement request, no typed prefix is pruned, and no context survives the command. There is no engine rerun, batch reassessment, or third cursor probe. The public Engine API, verification module, and REPL state contract remain unchanged. The current reference is docs/length-ranking.md; the command-authority predecessor, historical refill checkpoint, and private outcome seam are recorded in docs/reports/2026-08-15-command-authorized-length-filtering.md and docs/reports/2026-08-15-lane-local-length-survivor-refill.md, and docs/reports/2026-08-16-explicit-synthesis-lane-outcomes.md. The historical nominal-state foundation is recorded in docs/reports/2026-08-16-nominal-length-counterexample-bank-state.md; its runtime successor is docs/reports/2026-08-16-filter-only-length-counterexample-bank-context-runner.md. The command-local scheduler successor is docs/reports/2026-08-16-command-local-length-counterexample-bank-scheduler.md. The later Engine-only observation foundation is docs/reports/2026-08-16-opaque-detailed-synthesis-cursor-foundation.md. Its progressive Main runtime successor is docs/reports/2026-08-16-progressive-same-run-length-filter-batching.md. Dated reports under docs/reports/ are historical engineering records rather than current file contracts. When assessment is disabled, an ordinary or explicit no-contract rank remains the exact identity

and leaves no trace; explicit `filter` and every contract form instead return their pre-IO request errors.

When live work is required, the fixed `execve-check` launcher inherits Djex’s current descriptor ownership rule. Resource-producing worker results are published masked through one terminal completion cell; every native descriptor strategy shares one rollback-protected raw-child handoff; and interruptible initialization resumes only after the opaque process has accepted cleanup ownership. The child `exec-status` handle is closed on every normal, synchronous-failure, deadline, and asynchronous-interruption exit. These rules prevent leaked children and handles; they do not strengthen executable attestation, solver status, or behavioral evidence.

The pure searches answer in microseconds to milliseconds; the cost center is backend verification, a few hundred milliseconds per candidate. A wall-clock guard (default 20 s; `:set synth-timeout N` changes it for the session and `LEANT_SYNTHTIMEOUT` seeds it at startup) covers the quantified goals whose bounded instantiation can widen the space. Baseline and provider work share one absolute command deadline. Filter mode carries it through both batches and both classical routes; rank and disabled modes instead give each reached classical route a fresh configured duration. Zero waits indefinitely, and hitting a finite boundary is reported as “no answer,” never as a verdict.

5.11 What `:synth` will not do

Some plain statements of present limits. It will not synthesize recursion or induction. Complete compatible recursive families contribute bounded positive construction to Djinn and one constructor layer of elimination to Exference; recursive fields are never opened again. When the usual all-inputs-used Exference search has no candidate, Leant retries the same query with omissions allowed, which can project an impredicative payload while rendering the unopened recursive tail as `_`. It will not analyze dependent structure; atoms are transported, never opened. Opaque and bound-variable proper-type applications retain their ordered arguments, and qualifying inductives share one exact parametric family across distinct `Option/List` occurrences. Term-indexed, dependent, ambiguous, or incompatible families keep their conservative fallback. Quantifier handling beyond the propositional fragment follows Djex’s bounded rank-N and guarded impredicativity rules, which are sound and terminating over an undecidable space and widen as Djex improves. The current plan family is exhaustive through eleven independent quantified sites; a twelve-site six-open/six-opaque goal is the next deliberate occurrence-plan gap, while hypothesis instantiation chains with seven or more binders remain opaque. Such higher-rank goals answer “no term found within bounds” and nothing stronger. And although Djex can resolve type-class evidence itself, Leant leaves instance arguments to Lean’s elaborator during verification, which sits nearer the instance tables anyway. A non-dependent instance-implicit goal binder is retained as a render-only slot: it contributes no dictionary premise to either engine, but Leant inserts its wildcard before later lambda arguments and withholds complete negative evidence across the erased boundary. Constrained rank-N hypotheses therefore keep their instance applications implicit without shifting the candidate’s term binders. Provider discovery now consults active Lean instance heads only to obtain bounded ground-kinded assignments for the exact provider whose erased constraints they collectively solve. Constraint-only vacuous higher-kinded

binders are supported when each argument kind is a bounded `Type -> ... -> Type` chain. Canonical nominal `Prod` and `Sum` are supported bare or partially applied when their total arity is two, with direct occurrences mapped to the boxed-pair and `Either` identities. Legacy structural payloads and `And`, `PProd`, `Or`, `PSum`, `Iff`, and `Not` remain excluded. Leant never serializes dictionaries or general instance resolution into `Djex`. A selected head is accepted only after its own subgoals and every remaining provider constraint close in one isolated metavariable context. Up to 32 heads are inspected and 16 distinct complete vectors retained per provider, with at most 32 vectors passed in total. A closed contextual assignment such as `{a : Type} → [Inhabited a] → a → a` may cross this exact-provider lane only when serialization records the nominal class application and its ordered bounded-kinded arguments as `FExactContext`. Ground higher-kinded arguments such as `List` retain their canonical nominal head and residual kind arity. In a supported provider-source context, a lexically enclosing `FAll` variable may also occur bare or partially applied only to proper-type arguments. `Djex` checks the reconstructed class-parameter kind and rendering restores the exact Lean class head. The same contextual path supports canonical `Prod` and `Sum`; the direct path preserves their structural `Djex` identities. `Open`, `incomplete`, `wrong-arity`, `depth-truncated`, `malformed`, `free`, `dictionary-dependent`, `term-indexed`, and legacy raw `FInst` vectors fail closed. Ordinary goals retain render-only `FInst` slots. `Exference` retains supported provider-source `FExactContext` nodes on its plain, binder-only, and accepted-fact lanes; `Djinn` and a no-group exact-evidence fallback erase them. Accepted contextual vectors contribute replay-isolated ground facts rather than entering the context-free assignment adapter. This provider-local channel can expose a polytype absent from the query, but it is bounded checked evidence rather than general impredicative inference. The fragment is bounded on purpose: inside it, answers are fast and verified. Refutations are complete for the provider-free structural calculus but provisional with respect to the live environment. Leant searches the bounded, best-effort provider inventory before reporting one: a verified provider candidate overrides it, while an empty or unsuccessful provider search preserves it as the original sound fallback. The resulting verdict is not an exhaustive theorem about every declaration or axiom in the Lean environment.

6 How it works, briefly

A `:synth` query passes through five checked stages. A Lean metaprogram (compiled once into a cached side environment) elaborates the goal and serializes it into the engines' fragment: structural connectives are decomposed; qualifying inductives expand into constructor lists; proper-type applications retain rigid or bound heads and ordered arguments; everything else becomes an α -normalized opaque atom. The fragment translator either accepts the result or refuses with a reason. The cached synthesis environment remembers the session history it has replayed: an unchanged history is reused, an append replays only its new suffix, and undo or another non-prefix change rebuilds from the cached import-and-serializer base. Generated result bindings participate in this replay so later goals can mention them without flushing the provider cache. The selected engine searches; its candidates are rendered back into Lean syntax, with constructor names restored, binders named by role, and anonymous-constructor and leading-dot spellings preferred over fully qualified fallbacks. Finally the backend re-elaborates each candidate against the original goal, and only survivors are shown and bound.

The engine boundary is narrow by design: the searches behind it are swappable, and nothing they emit is believed until Lean has type-checked it.

An unpickled environment is an explicit history barrier. Leant copies it to a private process-lifetime artifact so backend restart can restore that base before replaying the commands entered afterward; the reconstructed undo stack therefore still bottoms out at the snapshot. `:reset` and `:load` leave snapshot mode, whereas `:import` requires that explicit transition because imports cannot be added to an opaque Lean environment. A Leant `:pickle` also writes a content-fingerprinted, serializer-ABI-keyed synthesis companion. Foreign or stale sidecars are ignored without rejecting the main upstream snapshot; if that snapshot exposes Lean's metaprogramming API, Leant compiles current synthesis tooling directly over it.

7 Status

Leant is a working prototype under active development. The synthesis design and its phased history are documented in `docs/SYNTHESIS_PROPOSAL.md`; transcript-level golden tests live in `test/`. The project began inside the *ProveIt* repository and was split out with its history intact. Expect sharp edges: the project exists to find out how far a conversational, verification-gated synthesis workflow can be pushed, and it moves at the speed of that question.